Data Analysis with Topic Models for Communications Researchers

1   Abstract

We present a non-technical introduction to data analysis with topic models for communications researchers. We motivate the discussion with a research question from social media communications research. We then discuss statistical aspects of topic models as we illustrate these methods with data from The New York Times. We complement our discussion with computer code (in the R computing language) that implements our methods. We close with thoughts about the future value of topic modeling to communications researchers.

2   Motivation

We are in the big data era. Social media users inundate us with status updates and tweets, while bloggers share their views on current events. These new media interact with more established media, such as print news media, television, and radio. What do they have in common? One answer is that they all can be viewed as data sources in which words - whether written or spoken, tweeted or blogged - are the data.

Think about how we currently browse this data. Let's suppose that we want to read about mass shootings. We might search for the key words "mass" and "shooting". Doing so, we will find some articles that deal with mass shootings, but other results may deal with other types of "shootings", such as film shooting or accidental shooting. Even if we specify that the two-word phrase "mass shooting" must appear in our results, we will find some articles

1

that merely mention mass shootings, and we'll miss relevant articles that don't explicitly use the term "mass shooting".

D. M. Blei (2012) envisions an alternative search strategy in which we search for articles that are *about* "mass shootings". Some of the resulting articles would have the phrase "mass shooting", but containing that two-word phrase is not required; the sole requirement is that the article focus on mass shootings. That is, we'd like to see all of the articles that have the theme, or "topic", mass shooting. Topic modeling is one method that may ultimately enable us to do the theme-based browsing that D. M. Blei (2012) suggests.

"Topic modeling" refers to a collection of statistical and machine learning methods that have the goal of discovering latent topics in a large collection, or "corpus", of texts. The tremendous rise in computing speed and memory capacity, coupled with the increasing availability of digitized texts, has enabled researchers working at the interface of quantitative methods and social sciences to treat written texts as data.

While data analysis of documents is still in its infancy, scientists nevertheless have made great progress towards computational dissection and interpretation of texts. We detail below, with limited use of statistical terminology, how these methods work and why they may be useful in communications research. We illustrate topic modeling methods in the analysis of all New York Times articles from three days in March 2016. We also provide an appendix with computer code (in the R computer programming language) that implements these methods.

3   Background

LDA is an extension of an earlier text analysis method called "probabilistic latent semantic analysis" (pLSA) (Hofmann, 1999). Hofmann (1999) described pLSA as a "latent class model for factor analysis of count data". The novelty of LDA, compared to pLSA, is that LDA places a Dirichlet prior on the probability distributions of words. This use of a Dirichlet prior makes

computations convenient because of a Bayesian statistical concept called "conjugacy". That is, the posterior probability distribution, which is the object of statistical inference, is also a Dirichlet distribution.

In the 13 years since the publication of D. M. Blei, Ng, & Jordan (2003), researchers have developed a wide variety of extensions to LDA. Many of these extensions of LDA accommodate relationships among topics. Widely used extensions of LDA include correlated topic models (D. M. Blei & Lafferty, 2007), hierarchical topic models (D. Griffiths & Tenenbaum, 2004), author-topic models (Rosen-Zvi, Griffiths, Steyvers, & Smyth, 2004), and dynamic topic models (D. M. Blei & Lafferty, 2006). Looking towards the future, researchers will need to fine tune and extend existing methods to accommodate new text data structures. One current research frontier is in topic modeling of tweets on Twitter. Twitter data present challenges because they are streaming and each message is no more than 140 characters.

Topic models have found applications in a wide range of disciplines. Pritchard, Stephens, & Donnelly (2000) used a statistical model much like LDA with genetic marker data to infer population structure and relatedness in human subjects. DeRose, Roy, & Boehm (2014) and Roy, DeRose, & Boehm (2013) used LDA (and other methods) in efforts to characterize themes from a collection of Victorian novels. Other examples include a 2008 study by Hall, Jurafsky, & Manning (2008), in which they applied LDA to examine the history of ideas. Grimmer (2010) applied a hierarchical topic model to study agendas in Senate press releases. We expect that the importance of topic modeling methods will continue to grow as digitized texts and digital media become more abundant.

4 Box's Loop & Data analysis

The perspective that guides our data analysis is based on ideas of the University of Wisconsin-Madison statistician George Box and coworkers [box1962useful]. D. M. Blei (2014) uses the ideas of Box and colleagues in the specific case of iterative refinement of topic models. Box

(1976) articulated a process of scientific inquiry and statistical collaboration in which scientists put forth a hypothesis about the natural world, and, with assistance and guidance from statisticians, design experiments to test the scientific hypothesis. After statistical analysis of the experimental data, the scientific hypothesis is refined, which leads to design and implementation of another experiment, and the iterative process between scientific hypothesis and experiment continues. Box (1976) also suggested that a statistician, working in collaboration with scientists, might use scientific questions as motivation to develop new statistical methods in experimental design and data analysis. In this sense, there is a second iterative loop by which a statistical researcher develops novel statistical methods because of their immediate need to answer a scientific research question.

D. M. Blei (2014) adapts these iterative processes to the case of latent variable models, such as LDA and other topic models. He argues that we view the use of a topic model for a specific data analysis task as an iterative procedure in which one proposes a simple topic model for the data analysis, fits the model, interprets the results, and then, if needed, refines the topic model with the goal of achieving data analysis results that are more consistent with the research goals.

A critical step in this process is that of model critiquing. D. M. Blei (2014) suggests using posterior predictive checks, evaluating performance of desired tasks, and considering prediction performance in this step. Recent research (Mimno & Blei, 2011) has made progress in developing methods for this last step in the iterative process articulated by D. M. Blei (2014).

5   Latent Dirichlet Allocation

D. M. Blei et al. (2003) introduced LDA as a (generative) statistical model in 2003. Although others had described similar statistical models (Pritchard et al., 2000), D. M. Blei et al. (2003) first applied the statistical model to text analysis. As we noted above, researchers

in a wide range of disciplines - including genetics, linguistics, psychology, political science, and others - have enthusiastically adopted LDA and related models to further their research efforts.

As D. M. Blei (2012) writes, the key to understanding LDA is to recognize that a given document - be it a research article, a novel, or a blog post - exhibits multiple topics. Each topic, in turn, is, in a technical sense, a probability distribution over words. For example, a topic related to evolution may heavily weight the words "evolution", "evolutionary", "biology", "phylogenetic", and "species". In a given collection of documents, which we term a "corpus" of documents, we assume that relatively few topics - on the order of 10 to 50 for most analyses - are present.

LDA models have a hiearchical structure in which words make up documents, and a collection of documents is a corpus. The corpus is assumed to have (unobserved) topics, or themes. The purposes of LDA, then, are to discover the unobserved topics from the texts and to characterize documents by the topics that they contain.

5.1   LDA Assumptions

One key assumption is often called the "bag of words" assumption. It states that the order of words (and the order of topics) in documents isn't important. Rather, we think of the process of generating words in a document as a matter of first choosing a probability distribution over topics, then, for each word in the document, using the selected probability distribution over topics to choose a topic, and then choosing a word from the corresponding probability distribution over the vocabulary. Hence, we allow for documents to reflect multiple topics and for each document to contain topics distinct proportions. Those familiar with probability theory will recognize the "bag of words" assumption as a statement of what probabilists call "exchangeability".

Hidden structures play key roles in LDA. While documents and their words are observed,

there remain three levels of hidden structures. The topics, the per-document topic distributions, and the per-document, per-word topic assignments constitute the unobserved variables that are the goal of our inference procedures. In other words, our statistical challenge, when working with topic models, is to use the observed structures - documents and their words - to infer the three types of hidden structures.

LDA's flexibility and adaptability have reduced the impact of its initial limitations. In the last decade, researchers have devised extensions of LDA, such as "dynamic topic models", that enable one to model topic evolution over time(D. M. Blei & Lafferty, 2006). Other extensions of LDA attempt to account for correlations among topics (D. M. Blei & Lafferty, 2007, Li & McCallum (2006)).

6 LDA with print media

Fitting LDA and related models of print media enables the user to discover topics. Let's suppose that we want to understand the themes that appear in a collection of print media articles. One strategy is to read every article in the collection. We could then assign one or more themes to each article and, at the end of our analysis, we would have a collection of themes. However, an alternative approach, in keeping with the theme-browsing ideas of D. M. Blei (2012), is to use LDA to identify topics, which we treat as themes, in the collection of articles.

To demonstrate our topic modeling approach to discovering themes, or topics, in print media, we used the Lexis Nexis Academic database to acquire 746 New York Times (print) articles from March 19, 20, and 21 of 2016.

We loaded the raw downloaded text files into a computer program called R, then partitioned the files into the 746 articles. To each article, we applied a series of text processing steps. We split each article into its distinct words and removed white space and punctuation before

converting all text to lower-case. At this stage, we had a collection of words for each of the 746 distinct articles. After removing words that occurred fewer than three times in the collection of articles and those that appeared on our "stop list" of syntactical short words, including articles and prepositions. With the remaining words, we created the document-term matrix. The document-term matrix is much like an Excel spreadsheet, where each row represents a document from the corpus and each column represents a word from the vocabulary. The cells in the matrix contain the number of times that each word appears in each document. Entering the document-term matrix into a computer program (R), we fit an LDA model with 20 topics.

7   LDA with social media

We downloaded tweets from the Twitter streaming API on the same three consecutive days (March 19, 20, 21 in 2016). Due to the large volume of downloaded tweets, we limited study to a one-hour period on each day (approximately 12pm Central time to 1pm Central time). The downloaded tweets constitute approximately one percent of all tweets during the selected time periods, according to Twitter's documentation. We then removed tweets that were in languages other than English, which left us with a total of 193211 tweets.

At this point, we made an important modeling decision in which we decided to treat each tweet as an independent document. This was an important decision because tweets are forced, by Twitter, to be short (no more than 140 characters in length). Because of their short length, it may be incomplete to think of a single tweet as a full document. Yet, as we'll see in the Results section, the model fitting yielded interpretable topics.

Processing our tweets requires removal of punctuation, URLs, and graphical characters that appear in tweets. After completing that step, we divided each tweet into its words and fit LDA models. As with the print media analysis, we chose to fit topic models with 20 topics.

8   Inference in topic models

Existing strategies for fitting topic models can be divided into two classes: 1. variational Bayes methods and 2. sampling methods. Variational Bayes methods try to approximate the posterior probability distribution by using optimization strategies from computer science. Alternatively, sampling methods, such as those based on Markov chain monte carlo (MCMC) approaches, draw samples from (an approximation to) the posterior probability distribution and work with those samples to approximate model parameters.

Computational implementations of both classes of strategies are freely available. We include R computing code in the appendix.

9   Interpreting results of topic modeling

Statistical inference for topic models yields estimates for topics (i.e., probability distributions over words), assignments of words (within documents) to topics, and weights of topics in each document. However, the user must impose meaning on the topics. For instance, a topic may put heavy weights on the words "genetics", "gene", "regulation", "DNA", "transcription", and "RNA". The data analyst needs to identify the similarities among these words - namely, that they describe concepts related to genetics and molecular biology. Fortunately, topic modeling procedures often discover topics that can be summarized in one or two words once the data analyst has inspected the most heavily weighted terms (for a given topic).

10   Visualizing topic models

Visualizing topic modeling results is an active area of research. We present below strategies that involve static and interactive displays. Word clouds are a widely used method for presenting topic modeling results. Unfortunately, the standard approach requires a distinct

word cloud for each topic. For models with more than ten topics, manually examining wordclouds becomes unwieldy.

Word cloud construction is straightforward and implemented with freely available software, such as the R package "wordcloud". For a single topic, the default method for creating a word cloud involves the identification of the most heavily weighted words in the topic. The user may choose to assign font sizes that are proportional to the weights of each word.

One newly developed method, which is implemented in the LDAvis R package (Sievert & Shirley, 2015), uses a singular value decomposition of the fitted document-topic matrix to calculate principal components. Each topic is then plotted in two dimensions, where the axes represent the first two principal components. The LDAvis R package enhances this second approach by making the figures interactive with D3 javascript.

For our analyses of tweets and New York Times articles, we created wordclouds for each topic in the resulting models.

11 Results from New York Times Articles Analysis

Examination of the resulting word clouds demonstrates that LDA is able to discover coherent themes from a text corpus. We see that our wordcloud in Figure 2 contains words, such as "ballet", "song", "music", and "film", that constitute the theme that we can label as "entertainment". The wordcloud in Figure 3 may be less clearly exhibiting a coherent theme, but it does contain a number of words, such as "boat", "river", "expedition", related to travel & adventure. Figure 4 is difficult to interpret. On the one hand, it contains "books", "review", and "story", which might suggest a literature theme, but at the same time it contains a collection of numbers. Figure 5 forms a cohesive collection of words about sports.

12   Results from Tweets Analysis

13   Discussion

We have presented preliminary results of our analysis of New York Times articles and tweets from March 19, 20, and 21, 2016. Looking back at Figure 1, we see that we've omitted the model critiquing step (D. M. Blei, 2014). Despite this, we still managed to discover meaningful topics in both the tweets and the New York Times articles with our initial models. Development of model critiquing strategies is one of our immediate research goals.

14   Future directions

Analysis of new data sources and structures often requires development of new statistical and computational methods. Topic models, despite their newness, have already contributed to scholarly pursuits in a wide range of social science, science, and humanities disciplines. Topic model-based analyses in journalism and communications research provide a practical method for summarizing and exploring the vast array of textual data that we encounter. Furthermore, analysis with topic models complements close reading and human annotation of texts. To fully realize the potential of topic modeling in communications research, we need to characterize the extent to which manual annotation of texts coincides with topic model-based annotation of texts. When such foundations are laid, we will realize a useful version of a topic-based browser for texts, as D. M. Blei (2012) envisions.

15   Online resources

David Mimno, a Cornell University scholar, curates an annotated bibliography of topic modeling research(Mimno, 2016). His bibliography is available at this url: http://mimno.

infosci.cornell.edu/topics.html

16   Computational implementation of LDA with R

We present below code for using LDA (with a collapsed Gibbs sampler) in the R statistical computing language (R Core Team, 2015).

```r
library(knitr)
opts_chunk$set(echo = FALSE, message = FALSE, warning = FALSE,
    cache = TRUE, tidy = TRUE, tidy.opts = list(blank = FALSE,
        width.cutoff = 60))  # hide source code in the document
tx1 <- scan(file = "data/The_New_York_Times2016-03-21_17-02.TXT",
    what = "character", blank.lines.skip = TRUE, sep = "\n",
    encoding = "UTF-8", skipNul = TRUE)
tx2 <- scan(file = "data/The_New_York_Times2016-03-21_17-04.TXT",
    what = "character", blank.lines.skip = TRUE, sep = "\n",
    encoding = "UTF-8", skipNul = TRUE)
tx <- c(tx1, tx2)
library(wordtools)  #load R package 'wordtools'
tx_list <- split_tx(tx = tx, patt = "Copyright 20")
library(stringr)
tx_list2 <- sapply(FUN = function(x) x[-(1:which(str_detect(string = x,
    pattern = "LENGTH")))], X = tx_list)
myfun <- function(x, pattern = "Web Blog") {
    collapsed <- paste(x, collapse = " ")
    !stringr::str_detect(collapsed, pattern = pattern)
}
```

```r
indswb <- sapply(FUN = myfun, X = tx_list2)

indsurl <- sapply(FUN = myfun, pattern = "URL", X = tx_list2)

library(magrittr)

good_art <- tx_list2[indswb] %>% sapply(FUN = function(x) {

    x[-(which(str_detect(x, "URL")):length(x))]

})

library(tm)

stopwords <- c(tm::stopwords("SMART"), "dr", "mr", "ms", "mrs")  # add titles to stopli

good2 <- sapply(FUN = function(x) paste(x, collapse = " "), X = good_art) %>%

    stringr::str_split(pattern = " ") %>% sapply(FUN = function(x) gsub("'",

    "", x)) %>% # remove apostrophes

sapply(FUN = function(x) gsub("[[:punct:]]", " ", x)) %>% # replace punctuation with sp

sapply(FUN = function(x) gsub("[[:cntrl:]]", " ", x)) %>% # replace control characters

sapply(FUN = function(x) gsub("^[[:space:]]+", "", x)) %>% # remove whitespace at begin

sapply(FUN = function(x) gsub("[[:space:]]+$", "", x)) %>% # remove whitesp Get rid of

sapply(FUN = function(x) str_replace_all(x, "#[a-z,A-Z]*", "")) %>%

    # Get rid of references to other screennames

sapply(FUN = function(x) str_replace_all(x, "@[a-z,A-Z]*", "")) %>%

    sapply(FUN = function(x) tolower(x)) %>% sapply(FUN = function(x) x[!(x ==

    "")]) %>% # remove elements that are ''

sapply(FUN = function(x) x[!(x %in% stopwords)])

# remove stopwords from

# http://cpsievert.github.io/LDAvis/reviews/reviews.html

# compute the table of terms:

n_min <- 3

term_table <- table(unlist(good2)) %>% sort(decreasing = TRUE)

term_table <- term_table[term_table >= n_min]
```

```r
vocab <- names(term_table)

get_terms <- function(x) {

    index <- match(x, vocab)

    index <- index[!is.na(index)]

    rbind(as.integer(index - 1), as.integer(rep(1, length(index))))

}

documents <- lapply(good2, get_terms)

# Compute some statistics related to the data set:

D <- length(documents)  # number of documents

W <- length(vocab)  # number of terms in the vocab

doc.length <- sapply(documents, function(x) sum(x[2, ]))  # number of tokens per docume

N <- sum(doc.length)  # total number of tokens in the data

term.frequency <- as.integer(term_table)  # frequencies of terms in the corpus

# MCMC and model tuning parameters:

K <- 20

G <- 5000

alpha <- 0.02

eta <- 0.02

# Fit the model:

library(lda)

set.seed(2016 - 3 - 22)

fit1 <- lda.collapsed.gibbs.sampler(documents = documents, K = K,

    vocab = vocab, num.iterations = G, alpha = alpha, eta = eta,

    initial = NULL, burnin = 1000, compute.log.likelihood = TRUE)

library(wordcloud)

for (i in 1:K) {

    cloud.data <- sort(fit1$topics[i, ], decreasing = TRUE)[1:50]
```

```r
    wordcloud(names(cloud.data), freq = cloud.data, scale = c(3,

        0.1), min.freq = 1, rot.per = 0, random.order = FALSE,

        col = 1 + i%%4)

}

tweet_dir <- "data/tweets/"

fns <- dir(tweet_dir)

out <- character(0)

for (file in fns) {

    tmp <- read.csv(file.path(tweet_dir, file), stringsAsFactors = FALSE)

    out <- rbind(out, tmp)

}

tweets <- out[out$lang == "en", ]

tw_txt <- tweets$text

library(stringr)

tw_words <- stringr::str_split(tw_txt, pattern = " ") %>% sapply(FUN = function(x) x[!(x

    "")]) %>% # remove elements that are ''

sapply(FUN = function(x) gsub("&amp", "", x)) %>% sapply(FUN = function(x) gsub("(RT|via

    "", x)) %>% sapply(FUN = function(x) gsub("@\\w+", "", x)) %>%

    sapply(FUN = function(x) gsub("[[:punct:]]", "", x)) %>%

    sapply(FUN = function(x) gsub("[[:digit:]]", "", x)) %>%

    sapply(FUN = function(x) gsub("[^[:graph:]]", "", x)) %>%

    # is this right???

sapply(FUN = function(x) gsub("http\\w+", "", x)) %>% sapply(FUN = function(x) gsub("[ 

    "", x)) %>% sapply(FUN = function(x) gsub("^\\s+|\\s+$",

    "", x)) %>% sapply(FUN = function(x) str_replace_all(x, " ",

    ""))

tw_words <- sapply(X = tw_words, FUN = function(x) str_replace(x,
```

```r
    "RT @[a-z,A-Z]*: ", "")) %>% sapply(FUN = function(x) str_replace_all(x,

    "#[a-z,A-Z]*", "")) %>% sapply(FUN = function(x) str_replace_all(x,

    "@[a-z,A-Z]*", "")) %>% # remove whitesp
sapply(FUN = function(x) tolower(x)) %>% sapply(FUN = function(x) x[!(x %in%

    tm::stopwords("SMART"))])
# remove stopwords remove 'rt' and ''
tw_words <- sapply(X = tw_words, FUN = function(x) x[!(x %in%

    c("", "rt"))])
# from http://cpsievert.github.io/LDAvis/reviews/reviews.html
# compute the table of terms:
n_min <- 3
term_table <- table(unlist(tw_words)) %>% sort(decreasing = TRUE)
term_table <- term_table[term_table >= n_min]
vocab <- names(term_table)
get_terms <- function(x) {
    index <- match(x, vocab)
    index <- index[!is.na(index)]
    rbind(as.integer(index - 1), as.integer(rep(1, length(index))))
}
documents <- lapply(tw_words, get_terms)
# Compute some statistics related to the data set:
D <- length(documents)  # number of documents
W <- length(vocab)  # number of terms in the vocab
doc.length <- sapply(documents, function(x) sum(x[2, ]))  # number of tokens per docume
N <- sum(doc.length)  # total number of tokens in the data
term.frequency <- as.integer(term_table)  # frequencies of terms in the corpus
# MCMC and model tuning parameters:
```

```r
K <- 20

G <- 5000

alpha <- 0.02

eta <- 0.02

# Fit the model:

library(lda)

set.seed(2016 - 3 - 22)

fit1 <- lda.collapsed.gibbs.sampler(documents = documents, K = K,

    vocab = vocab, num.iterations = G, alpha = alpha, eta = eta,

    initial = NULL, burnin = 1000, compute.log.likelihood = TRUE)

library(wordcloud)

for (i in 1:K) {

    cloud.data <- sort(fit1$topics[i, ], decreasing = TRUE)[1:50]

    wordcloud(names(cloud.data), freq = cloud.data, scale = c(3,

        0.1), min.freq = 1, rot.per = 0, random.order = FALSE,

        col = 1 + i%%4)

}
```

17   LDA with statistical terminology

We can also describe LDA with more formal statistical terminology. We let $\beta_1, \beta_2, ..., \beta_K$ be the K topics, where each β_k is a probability distribution over the vocabulary. Topic weights (or proportions) we denote by θ_d for the d^{th} document. We let $\theta_{d,k}$ be the topic proportion for topic k in document d. Topic assignments for the d^{th} document are z_d, where $z_{d,n}$ is the topic assignment for the n^{th} word in the d^{th} document. Observed words for the d^{th} document are w_d, where $w_{d,n}$ is the n^{th} word for the d^{th} document.

We then use the above notation to specify the joint probability distribution for all variables in the model:

$$p(\beta_{1:K}, \theta_{1:D}, z_{1:D}, w_{1:D}) = \prod_{i=1}^{K} p(\beta_i) \prod_{d=1}^{D} p(\theta_d) \left(\prod_{n=1}^{N} p(z_{d,n}|\theta_d) p(w_{d,n}|\beta_{1:K}, z_{d,n}) \right)$$

The vertical bars in the last expressions denote conditional probabilities. Note that the topic assignments $z_{d,n}$ depend on the document's topic proportions. Whatismore, the observed word $w_{d,n}$ has a probability distribution that depends on both the entire collection of topics, $\beta_{1:K}$, and the topic, $z_{d,n}$, assigned to that word.

References

Blei, D. M. (2012). Probabilistic topic models. *Communications of the ACM*, *55*(4), 77–84.

Blei, D. M. (2014). Build, compute, critique, repeat: Data analysis with latent variable models. *Annual Review of Statistics and Its Application*, *1*, 203–232.

Blei, D. M., & Lafferty, J. D. (2006). Dynamic topic models. In *Proceedings of the 23rd international conference on machine learning* (pp. 113–120). ACM.

Blei, D. M., & Lafferty, J. D. (2007). A correlated topic model of science. *The Annals of Applied Statistics*, 17–35.

Blei, D. M., Ng, A. Y., & Jordan, M. I. (2003). Latent dirichlet allocation. *The Journal of Machine Learning Research*, *3*, 993–1022.

Box, G. E. (1976). Science and statistics. *Journal of the American Statistical Association*, *71*(356), 791–799.

DeRose, C., Roy, C., & Boehm, F. (2014). Victorian eyes: Literary, statistical, and artistic perspectives on victorian novels – and dickens's unfinished murder mystery. *Significance*,

11(2), 40–43.

Griffiths, D., & Tenenbaum, M. (2004). Hierarchical topic models and the nested chinese restaurant process. *Advances in Neural Information Processing Systems*, *16*, 17.

Grimmer, J. (2010). A bayesian hierarchical topic model for political texts: Measuring expressed agendas in senate press releases. *Political Analysis*, *18*(1), 1–35.

Hall, D., Jurafsky, D., & Manning, C. D. (2008). Studying the history of ideas using topic models. In *Proceedings of the conference on empirical methods in natural language processing* (pp. 363–371). Association for Computational Linguistics.

Hofmann, T. (1999). Probabilistic latent semantic indexing. In *Proceedings of the 22nd annual international aCM sIGIR conference on research and development in information retrieval* (pp. 50–57). ACM.

Li, W., & McCallum, A. (2006). Pachinko allocation: DAG-structured mixture models of topic correlations. In *Proceedings of the 23rd international conference on machine learning* (pp. 577–584). ACM.

Mimno, D. (2016). Topic modeling bibliography. Retrieved from http://mimno.infosci.cornell.edu/topics.html

Mimno, D., & Blei, D. (2011). Bayesian checking for topic models. In *Proceedings of the conference on empirical methods in natural language processing* (pp. 227–237). Association for Computational Linguistics.

Pritchard, J. K., Stephens, M., & Donnelly, P. (2000). Inference of population structure using multilocus genotype data. *Genetics*, *155*(2), 945–959.

R Core Team. (2015). *R: A language and environment for statistical computing.* Vienna, Austria: R Foundation for Statistical Computing. Retrieved from https://www.R-project.org/

Rosen-Zvi, M., Griffiths, T., Steyvers, M., & Smyth, P. (2004). The author-topic model for

authors and documents. In *Proceedings of the 20th conference on uncertainty in artificial intelligence* (pp. 487–494). AUAI Press.

Roy, C., DeRose, C., & Boehm, F. (2013). Victorian eyes. Retrieved from http://victorianeyes.com

Sievert, C., & Shirley, K. (2015). *LDAvis: Interactive visualization of topic models.* Retrieved from https://CRAN.R-project.org/package=LDAvis

List of Figures

Figure 1: Iterative process for building, computing, critiquing, and applying topic models.

Figure 2: Wordcloud for one topic from a 20-topic model of 746 New York Times articles. Larger font size corresponds to greater weight of that word in this topic.

Figure 3: Wordcloud for one topic from a 20-topic model of 746 New York Times articles. Larger font size corresponds to greater weight of that word in this topic.

Figure 4: Wordcloud for one topic from a 20-topic model of 746 New York Times articles. Larger font size corresponds to greater weight of that word in this topic.

Figure 5: Wordcloud for one topic from a 20-topic model of 746 New York Times articles. Larger font size corresponds to greater weight of that word in this topic.

Figure 6: Wordcloud for one topic from a 20-topic model of 193,211 tweets. Larger font size corresponds to greater weight of that word in this topic.

Figure 7: Wordcloud for one topic from a 20-topic model of 193,211 tweets. Larger font size corresponds to greater weight of that word in this topic.

Figure 8: Wordcloud for one topic from a 20-topic model of 193,211 tweets. Larger font size corresponds to greater weight of that word in this topic.

Figure 9: Wordcloud for one topic from a 20-topic model of 193,211 tweets. Larger font size corresponds to greater weight of that word in this topic.